Conjugate Gradient Method

- direct and indirect methods

- positive definite linear systems

- Krylov sequence

- spectral analysis of Krylov sequence

- preconditioning

Prof. S. Boyd, EE364b, Stanford University

Three classes of methods for linear equations

methods to solve linear system $Ax = b$, $A \in \mathbf{R}^{n \times n}$

- **dense direct** (factor-solve methods)

  - runtime depends only on size; independent of data, structure, or sparsity
  - work well for n up to a few thousand

- **sparse direct** (factor-solve methods)

  - runtime depends on size, sparsity pattern; (almost) independent of data
  - can work well for n up to 10^4 or 10^5 (or more)
  - requires good heuristic for ordering

- **indirect** (iterative methods)

  - runtime depends on data, size, sparsity, required accuracy
  - requires tuning, preconditioning, . . .
  - good choice in many cases; only choice for $n = 10^6$ or larger

Symmetric positive definite linear systems

SPD system of equations

$$Ax = b, \qquad A \in \mathbf{R}^{n \times n}, \qquad A = A^T \succ 0$$

examples

- Newton/interior-point search direction: $\nabla^2 \phi(x) \Delta x = -\nabla \phi(x)$

- least-squares normal equations: $(A^T A)x = A^T b$

- regularized least-squares: $(A^T A + \mu I)x = A^T b$

- minimization of convex quadratic function $(1/2)x^T A x - b^T x$

- solving (discretized) elliptic PDE (*e.g.*, Poisson equation)

- analysis of resistor circuit: $Gv = i$

  - v is node voltage (vector), i is (given) source current
  - G is circuit conductance matrix

$$G_{ij} = \begin{cases} \text{total conductance incident on node } i & i = j \\ -(\text{conductance between nodes } i \text{ and } j) & i \neq j \end{cases}$$

CG overview

- proposed by Hestenes and Stiefel in 1952 (as direct method)

- solves SPD system $Ax = b$

  - in theory ($i.e.$, exact arithmetic) in n iterations
  - each iteration requires a few inner products in $\mathbf{R}^n$, and one matrix-vector multiply $z \to Az$

- for A dense, matrix-vector multiply $z \to Az$ costs n^2, so total cost is n^3, same as direct methods

- get advantage over dense if matrix-vector multiply is cheaper than n^2

- with roundoff error, CG can work poorly (or not at all)

- but for some A (and b), can get good approximate solution in $\ll n$ iterations

Solution and error

- $x^\star = A^{-1}b$ is solution

- $x^\star$ minimizes (convex function) $f(x) = (1/2)x^T Ax - b^T x$

- $\nabla f(x) = Ax - b$ is gradient of f

- with $f^\star = f(x^\star)$, we have

$$
\begin{aligned}
f(x) - f^\star &= (1/2)x^T Ax - b^T x - (1/2)x^{\star T} Ax^\star + b^T x^\star \\
&= (1/2)(x - x^\star)^T A(x - x^\star) \\
&= (1/2)\|x - x^\star\|_A^2
\end{aligned}
$$

i.e., $f(x) - f^\star$ is half of squared A-norm of error $x - x^\star$

- a relative measure (comparing x to 0):

$$\tau = \frac{f(x) - f^\star}{f(0) - f^\star} = \frac{\|x - x^\star\|_A^2}{\|x^\star\|_A^2}$$

(fraction of maximum possible reduction in f, compared to $x = 0$)

Residual

- $r = b - Ax$ is called the **residual** at x

- $r = -\nabla f(x) = A(x^\star - x)$

- in terms of r, we have

$$
\begin{aligned}
f(x) - f^\star &= (1/2)(x - x^\star)^T A(x - x^\star) \\
&= (1/2)r^T A^{-1} r \\
&= (1/2)\|r\|_{A^{-1}}^2
\end{aligned}
$$

- a commonly used measure of relative accuracy: $\eta = \|r\|/\|b\|$

- $\tau \le \kappa(A)\eta^2$ (η is easily computable from x; τ is not)

Krylov subspace

(a.k.a. controllability subspace)

$$
\begin{aligned}
\mathcal{K}_k &= \mathrm{span}\{b, Ab, \ldots, A^{k-1}b\} \\
&= \{p(A)b \mid p \text{ polynomial}, \ \deg p < k\}
\end{aligned}
$$

we define the *Krylov sequence* $x^{(1)}, x^{(2)}, \ldots$ as

$$
x^{(k)} = \underset{x \in \mathcal{K}_k}{\mathrm{argmin}} \ f(x) = \underset{x \in \mathcal{K}_k}{\mathrm{argmin}} \ \|x - x^\star\|_A^2
$$

the CG algorithm (among others) generates the Krylov sequence

Properties of Krylov sequence

- $f(x^{(k+1)}) \le f(x^{(k)})$ (but $\|r\|$ can increase)

- $x^{(n)} = x^\star$ (*i.e.*, $x^\star \in \mathcal{K}_n$ even when $\mathcal{K}_n \ne \mathbf{R}^n$)

- $x^{(k)} = p_k(A)b$, where p_k is a polynomial with $\deg p_k < k$

- less obvious: there is a *two-term recurrence*

$$x^{(k+1)} = x^{(k)} + \alpha_k r^{(k)} + \beta_k (x^{(k)} - x^{(k-1)})$$

for some α_k, β_k (basis of CG algorithm)

Cayley-Hamilton theorem

characteristic polynomial of A:

$$\chi(s) = \det(sI - A) = s^n + \alpha_1 s^{n-1} + \cdots + \alpha_n$$

by Caley-Hamilton theorem

$$\chi(A) = A^n + \alpha_1 A^{n-1} + \cdots + \alpha_n I = 0$$

and so

$$A^{-1} = -(1/\alpha_n) A^{n-1} - (\alpha_1/\alpha_n) A^{n-2} - \cdots - (\alpha_{n-1}/\alpha_n) I$$

in particular, we see that $x^\star = A^{-1} b \in \mathcal{K}_n$

Spectral analysis of Krylov sequence

- $A = Q\Lambda Q^T$, Q orthogonal, $\Lambda = \mathbf{diag}(\lambda_1, \dots, \lambda_n)$

- define $y = Q^T x$, $\bar{b} = Q^T b$, $y^\star = Q^T x^\star$

- in terms of y, we have

$$
\begin{aligned}
f(x) = \bar{f}(y) &= (1/2)x^T Q\Lambda Q^T x - b^T Q Q^T x \\
&= (1/2)y^T \Lambda y - \bar{b}^T y \\
&= \sum_{i=1}^{n} \left((1/2)\lambda_i y_i^2 - \bar{b}_i y_i \right)
\end{aligned}
$$

so $y_i^\star = \bar{b}_i/\lambda_i$, $f^\star = -(1/2)\sum_{i=1}^{n} \bar{b}_i^2/\lambda_i$

Krylov sequence in terms of y

$$y^{(k)} = \operatorname*{argmin}_{y \in \bar{\mathcal{K}}_k} \bar{f}(y), \qquad \bar{\mathcal{K}}_k = \operatorname{span}\{\bar{b}, \Lambda\bar{b}, \ldots, \Lambda^{k-1}\bar{b}\}$$

$$y_i^{(k)} = p_k(\lambda_i)\bar{b}_i, \qquad \deg p_k < k$$

$$p_k = \operatorname*{argmin}_{\deg p < k} \sum_{i=1}^{n} \bar{b}_i^2 \left((1/2)\lambda_i p(\lambda_i)^2 - p(\lambda_i)\right)$$

$$
\begin{aligned}
f(x^{(k)}) - f^\star \;&=\; \bar{f}(y^{(k)}) - f^\star \\[2mm]
&=\; \min_{\deg p < k} (1/2) \sum_{i=1}^{n} \bar{b}_i^2 \frac{(\lambda_i p(\lambda_i) - 1)^2}{\lambda_i} \\[2mm]
&=\; \min_{\deg p < k} (1/2) \sum_{i=1}^{n} \bar{y}_i^{\star 2} \lambda_i (\lambda_i p(\lambda_i) - 1)^2 \\[2mm]
&=\; \min_{\deg q \le k,\; q(0)=1} (1/2) \sum_{i=1}^{n} \bar{y}_i^{\star 2} \lambda_i q(\lambda_i)^2 \\[2mm]
&=\; \min_{\deg q \le k,\; q(0)=1} (1/2) \sum_{i=1}^{n} \bar{b}_i^2 \frac{q(\lambda_i)^2}{\lambda_i}
\end{aligned}
$$

$$\tau_k = \frac{\min_{\deg q \le k,\ q(0)=1} \sum_{i=1}^{n} \bar{y}_i^{\star 2} \lambda_i q(\lambda_i)^2}{\sum_{i=1}^{n} \bar{y}_i^{\star 2} \lambda_i}$$

$$\le \min_{\deg q \le k,\ q(0)=1} \left(\max_{i=1,\ldots,n} q(\lambda_i)^2 \right)$$

- if there is a polynomial q of degree k, with $q(0) = 1$, that is small on the spectrum of A, then $f(x^{(k)}) - f^\star$ is small

- if eigenvalues are clustered in k groups, then $y^{(k)}$ is a good approximate solution

- if solution $x^\star$ is approximately a linear combination of k eigenvectors of A, then $y^{(k)}$ is a good approximate solution

A bound on convergence rate

- taking q as Chebyshev polynomial of degree k, that is small on interval $[\lambda_{\min}, \lambda_{\max}]$, we get

$$\tau_k \leq \left(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \right)^k, \qquad \kappa = \lambda_{\max}/\lambda_{\min}$$

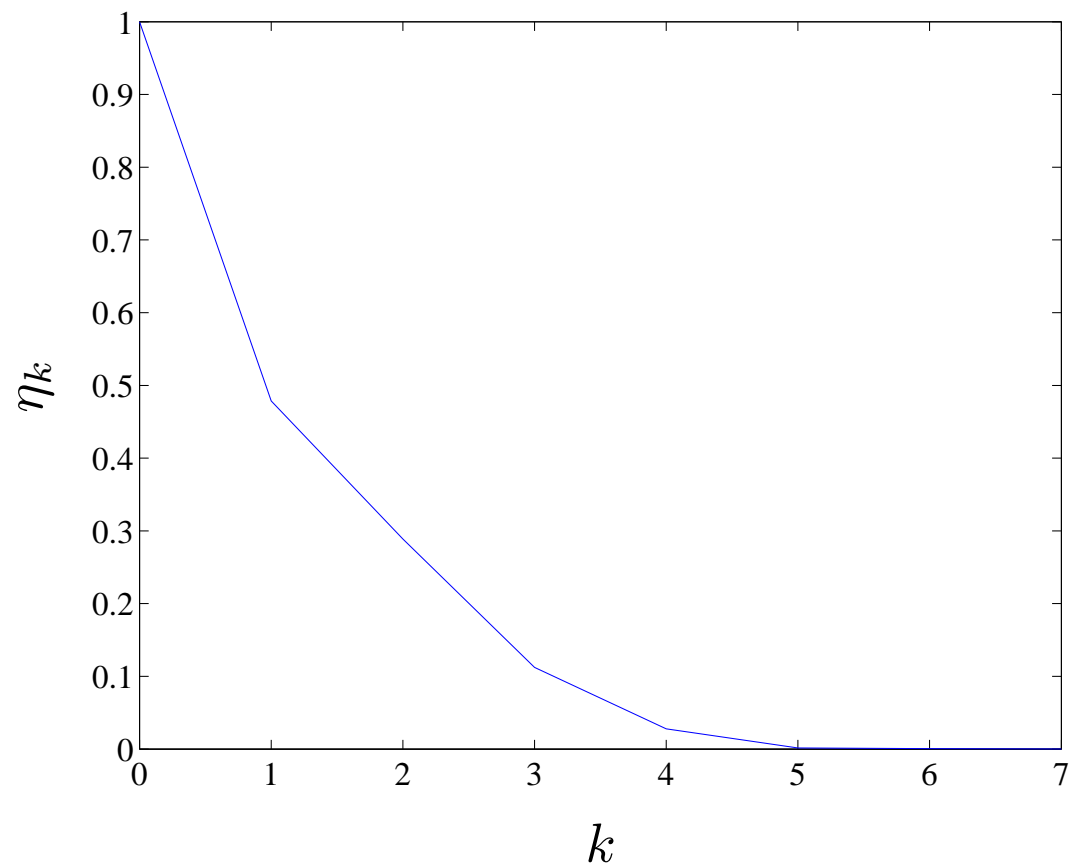- convergence can be much faster than this, if spectrum of A is spread but clustered

Small example

$A \in \mathbf{R}^{7 \times 7}$, spectrum shown as filled circles; p_1, p_2, p_3, p_4, and p_7 shown
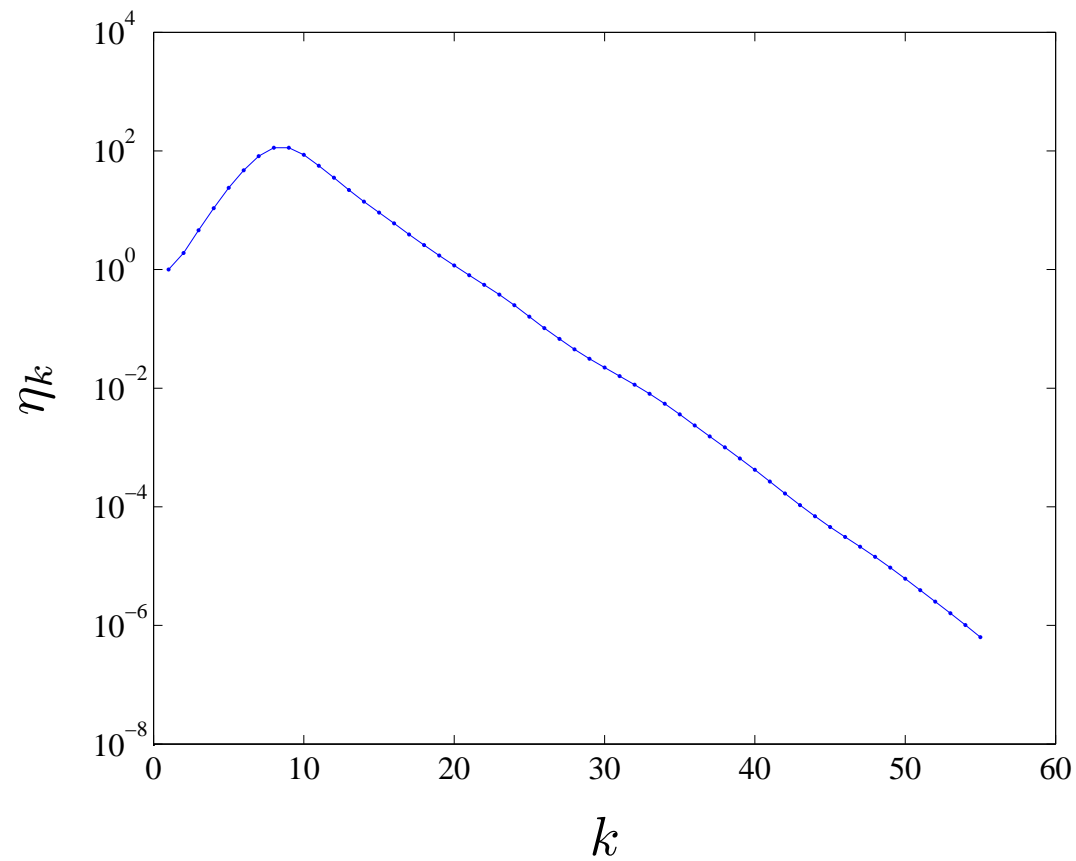
Convergence

Residual convergence

Larger example

- solve $Gv = i$, resistor network with 10^5 nodes

- average node degree 10; around 10^6 nonzeros in G

- random topology with one grounded node

- nonzero branch conductances uniform on $[0, 1]$

- external current i uniform on $[0, 1]$

- sparse Cholesky factorization of G requires too much memory

Residual convergence

CG algorithm

(follows C. T. Kelley)

$$x := 0, \quad r := b, \quad \rho_0 := \|r\|^2$$

for $k = 1, \ldots, N_{\max}$
      quit if $\sqrt{\rho_{k-1}} \leq \epsilon \|b\|$
      if $k = 1$ then $p := r$; else $p := r + (\rho_{k-1}/\rho_{k-2})p$
      $w := Ap$
      $\alpha := \rho_{k-1}/p^T w$
      $x := x + \alpha p$
      $r := r - \alpha w$
      $\rho_k := \|r\|^2$

Efficient matrix-vector multiply

- sparse A

- structured ($e.g.$, sparse) plus low rank

- products of easy-to-multiply matrices

- fast transforms (FFT, wavelet, . . .)

- inverses of lower/upper triangular (by forward/backward substitution)

- fast Gauss transform, for $A_{ij} = \exp(-\|v_i - v_j\|^2/\sigma^2)$ (via multipole)

Shifting

- suppose we have guess $\hat{x}$ of solution $x^\star$

- we can solve $Az = b - A\hat{x}$ using CG, then get $x^\star = \hat{x} + z$

- in this case $x^{(k)} = \hat{x} + z^{(k)} = \underset{x \in \hat{x} + \mathcal{K}_k}{\operatorname{argmin}} f(x)$

  ($\hat{x} + \mathcal{K}_k$ is called *shifted Krylov subspace*)

- same as initializing CG alg with $x := \hat{x}$, $r := b - Ax$

- good for 'warm start', *i.e.*, solving $Ax = b$ starting from a good initial guess (*e.g.*, the solution of another system $\tilde{A}x = \tilde{b}$, with $A \approx \tilde{A}$, $b \approx \tilde{b}$)

Preconditioned conjugate gradient algorithm

- idea: apply CG after linear change of coordinates $x = Ty$, $\det T \neq 0$

- use CG to solve $T^T A T y = T^T b$; then set $x^\star = T^{-1} y^\star$

- T or $M = T T^T$ is called *preconditioner*

- in naive implementation, each iteration requires multiplies by T and T^T (and A); also need to compute $x^\star = T^{-1} y^\star$ at end

- can re-arrange computation so each iteration requires one multiply by M (and A), and no final solve $x^\star = T^{-1} y^\star$

- called *preconditioned conjugate gradient* (PCG) algorithm

Choice of preconditioner

- if spectrum of $T^T A T$ (which is the same as the spectrum of MA) is clustered, PCG converges fast

- extreme case: $M = A^{-1}$

- trade-off between enhanced convergence, and extra cost of multiplication by M at each step

- goal is to find M that is cheap to multiply, and approximate inverse of A (or at least has a more clustered spectrum than A)

Some generic preconditioners

- diagonal: $M = \mathbf{diag}(1/A_{11}, \ldots, 1/A_{nn})$

- incomplete/approximate Cholesky factorization: use $M = \hat{A}^{-1}$, where $\hat{A} = \hat{L}\hat{L}^T$ is an approximation of A with cheap Cholesky factorization

  - compute Cholesky factorization of $\hat{A}$, $\hat{A} = \hat{L}\hat{L}^T$
  - at each iteration, compute $Mz = \hat{L}^{-T}\hat{L}^{-1}z$ via forward/backward substitution

- examples

  - $\hat{A}$ is central k-wide band of A
  - $\hat{L}$ obtained by sparse Cholesky factorization of A, ignoring small elements in A, or refusing to create excessive fill-in

Larger example

residual convergence with and without diagonal preconditioning

CG summary

- in theory (with exact arithmetic) converges to solution in n steps

  - the bad news: due to numerical round-off errors, can take more than n steps (or fail to converge)
  - the good news: with luck ($i.e.$, good spectrum of A), can get good approximate solution in $\ll n$ steps

- each step requires $z \rightarrow Az$ multiplication

  - can exploit a variety of structure in A
  - in many cases, never form or store the matrix A

- compared to direct (factor-solve) methods, CG is less reliable, data dependent; often requires good (problem-dependent) preconditioner

- but, when it works, can solve extremely large systems